

# Deep Learning

## Lecture 3

# Recap

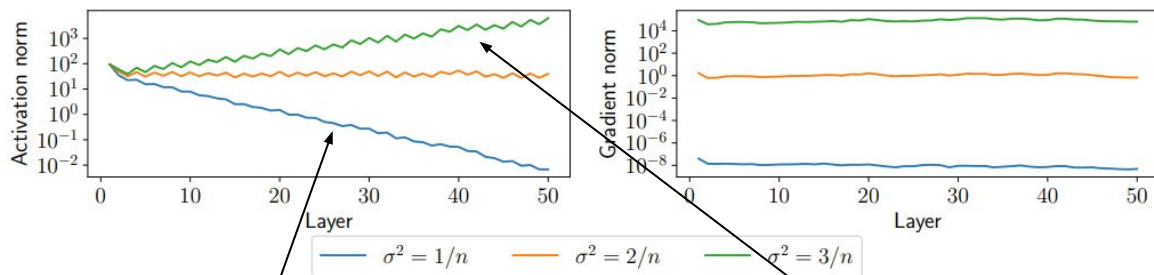
- Gradient descent for neural networks
  - SGD
  - Momentum
  - Adaptive learning rate
- Weight decay
  - L2 regularization in NN
  - Adam vs AdamW
- Dropout
  - Ensemble
  - Difference on `train()` and `eval()` modes

# Contents

- Vanishing and exploding gradient
- Weights initialization
- Internal Covariate Shift
- Normalization
  - Batch Norm
  - Layer Norm
  - Instance Norm
  - Group Norm
- Convolutional NN

# Motivation for weight initialization

- Suppose we have the following initialization:  $w_i \sim \mathcal{N}\left(0, \frac{c}{n}\right)$
- Calculate 1) Activation norm, and 2) Gradient norm for each layer



Vanishing Gradient

Exploding Gradient

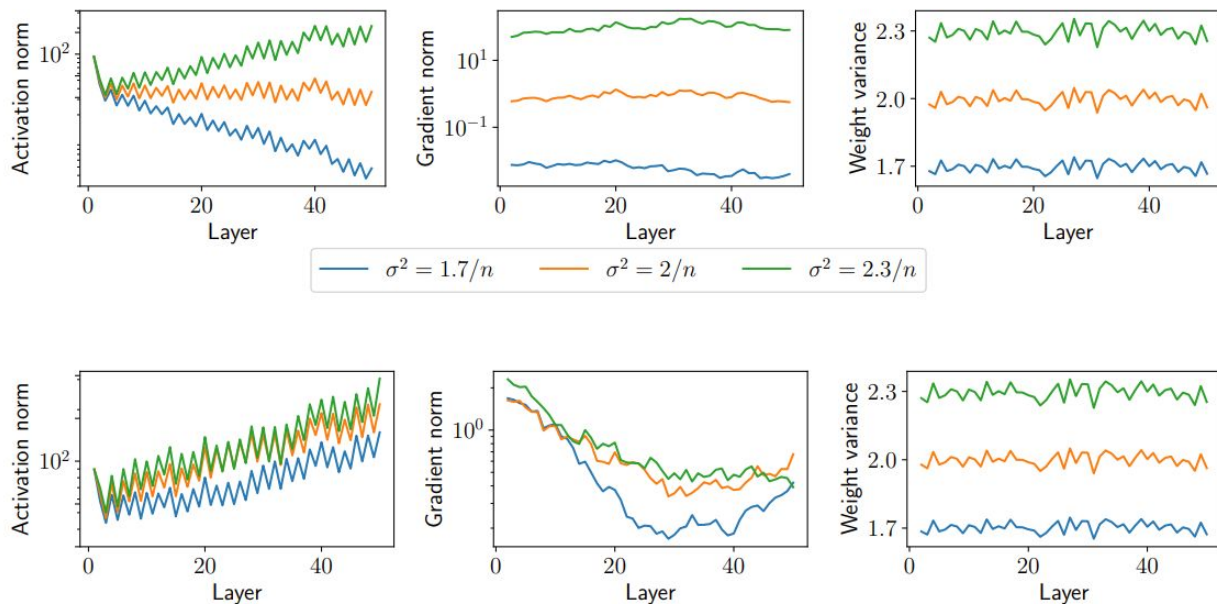
$\sigma^2 = 3/n \Rightarrow \text{NaN}$

$\sigma^2 = 2/n \Rightarrow \text{Works}$

$\sigma^2 = 1/n \Rightarrow \text{No progress}$

# Weights after training

The problem is even more fundamental, however: even when trained successfully, the effects/scales present at initialization **persist** throughout training



Train to 5%  
error on  
MNIST

# Weights Initialization

What kind of initialization can work?

- Zero initialization or constant initialization?\*
  - Degeneration of a neural network into a single neuron
- Random initialization?
  - Does not ensure the avoidance of vanishing and exploding gradient
- More smart approaches?
  - Formalize the required neural network properties

\*Why is it a bad idea? See [this post](#)

# Idea

To prevent the gradients of the network's activations from vanishing or exploding, we will stick to the following rules of thumb:

1. The *mean* of the activations should be zero.
2. The *variance* of the activations should stay the same across every layer.

$$a^{l-1} = g^{l-1}(z^{l-1})$$

$$z^l = W^l a^{l-1} + b^l$$

$$a^l = g^l(z^l)$$

Desired condition:

$$\mathbb{E}(a^{l-1}) = \mathbb{E}(a^l)$$

$$\text{Var}(a^{l-1}) = \text{Var}(a^l)$$

# Xavier and He initialization

Xavier initialization (for sigmoid and tanh)

$$W^{[l]} \sim \mathcal{N}\left(\mu = 0, \sigma^2 = \frac{1}{n^{[l-1]}}\right) \quad b^{[l]} = 0$$

He initialization (for ReLU)

$$W^{[l]} \sim \mathcal{N}\left(\mu = 0, \sigma^2 = \frac{2}{n^{[l-1]}}\right) \quad b^{[l]} = 0$$

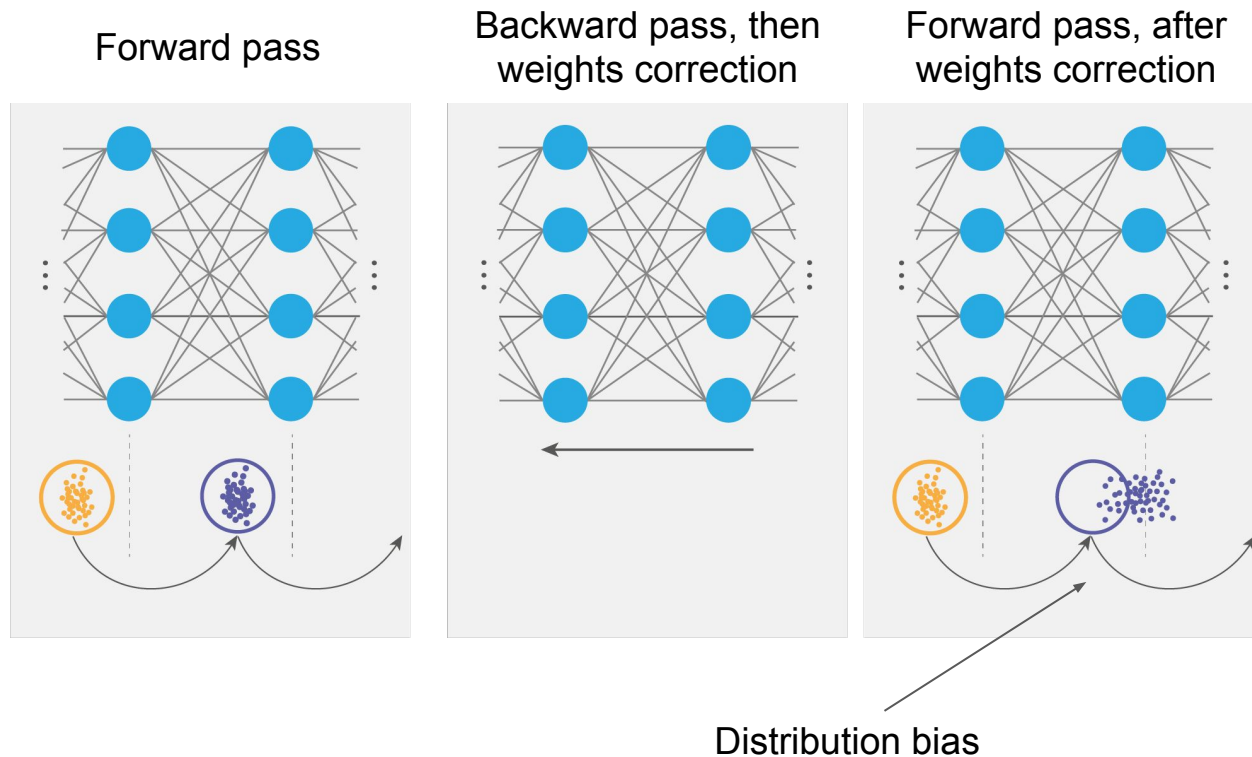
$W^{[l]}$  – weights

$b^{[l]}$  – bias

$n^{[l-1]}$  – hidden size



# Internal Covariate Shift

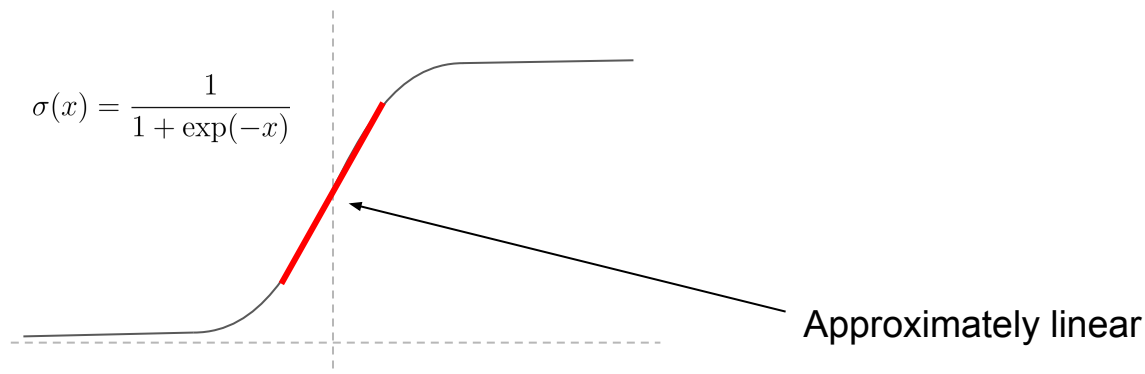


# Naive approach

- Simply normalize each feature (each neuron output), using the mean and variance from batch:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sigma_B + \varepsilon}$$

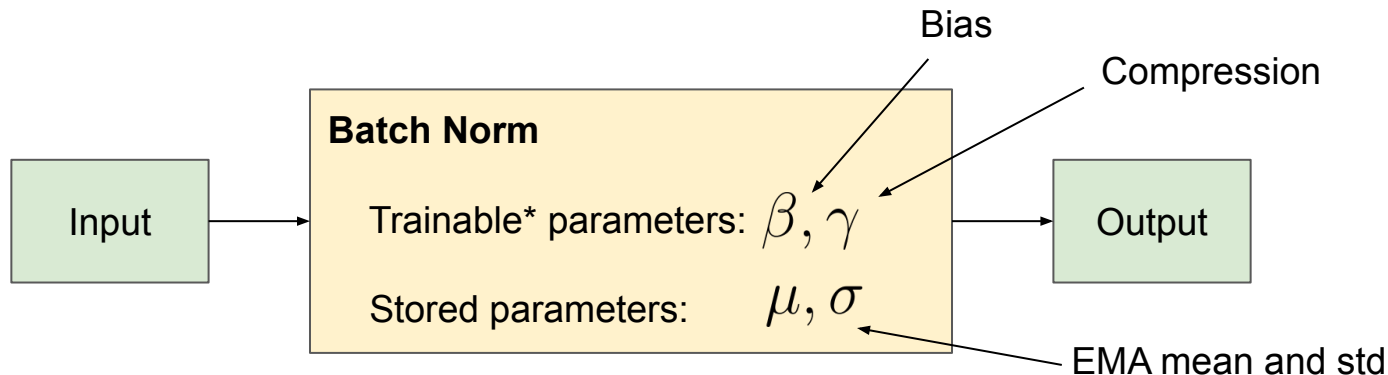
- After normalizing, an activation function seems to be linear



- Thus, we obtain just a linear one-layer network...

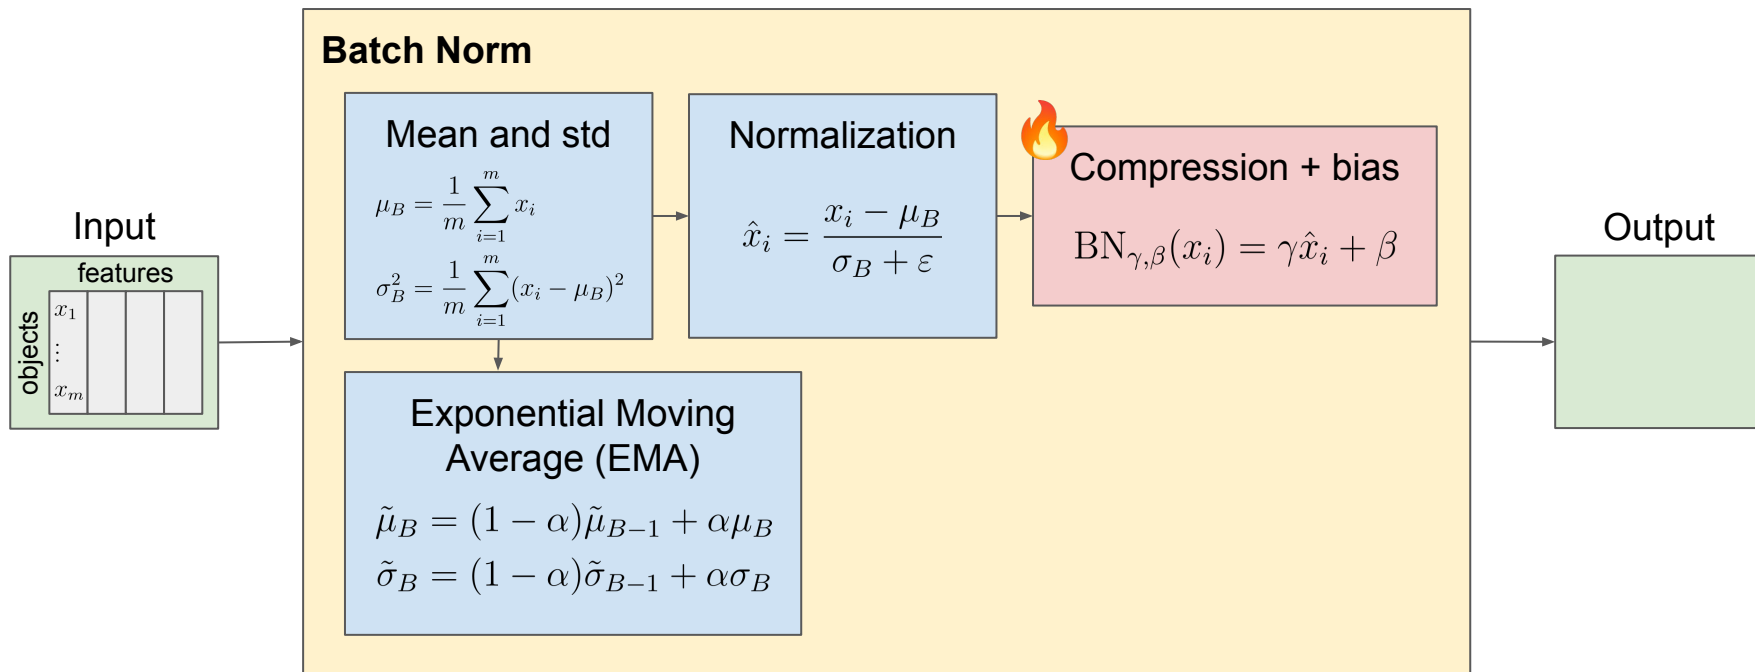
# Batch Normalization

- We can give the neural network the ability to move the distribution of layer outputs:
  - Bias
  - Compression
- During the **training stage**, we calculate statistics on the batch
- How to make a prediction on a single object from **test**?
  - Calculate mean and variance on the entire dataset? Not enough memory...
  - Estimate them? Exponential Moving Average (EMA)!



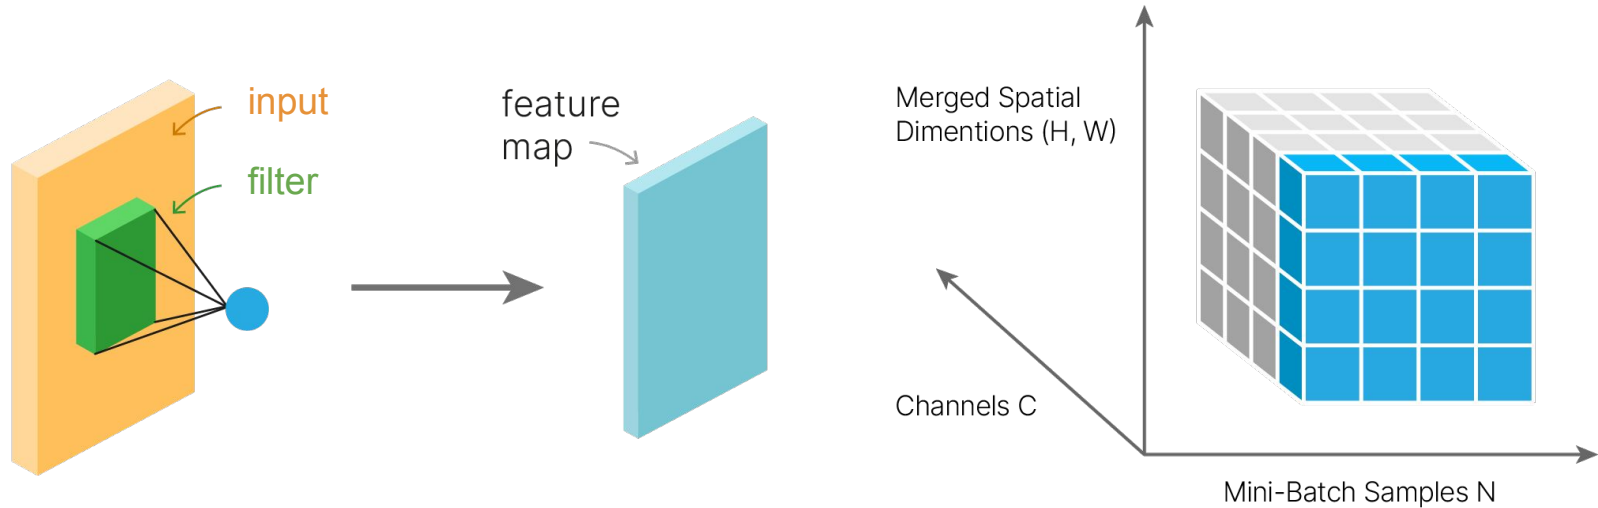
\*How to calculate the gradient? See [this post](#) on Kevin Zakka's Blog

# Batch Normalization: scheme



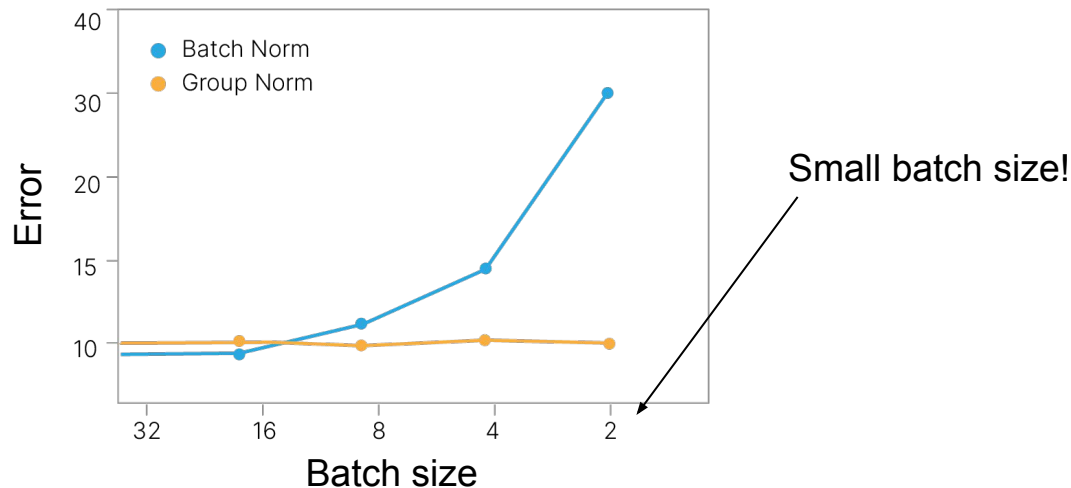
# Batch Normalization: visualization

- In case of the feature maps, Batch Norm is applied **for each channel** (like feature)



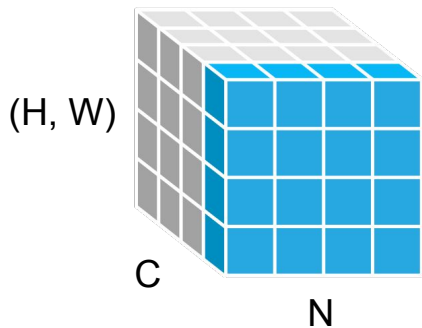
# Batch Normalization: Practical Recommendations

- Shuffle objects between epochs when using Batch Norm to ensure diverse batches for training parameters
- Remove bias in the layer following Batch Norm, as the bias parameter in Batch Norm takes on this role.
- Use before activation function
- **Smaller batch sizes result in poorer performance of Batch Norm**



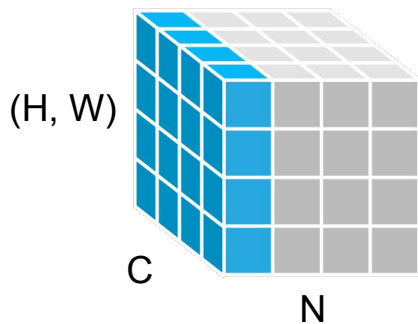
# Other Normalizations

Batch Norm



Classical architectures

Layer Norm



Modern architectures  
(e.g. transformers)

Instance Norm

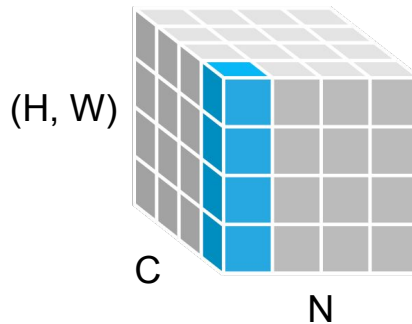
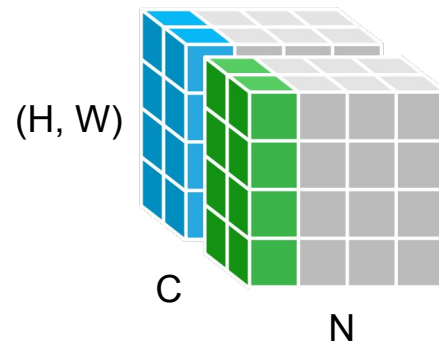


Image generation  
(contrast balance)

Group Norm



Used sometimes...

# Convolutional NN



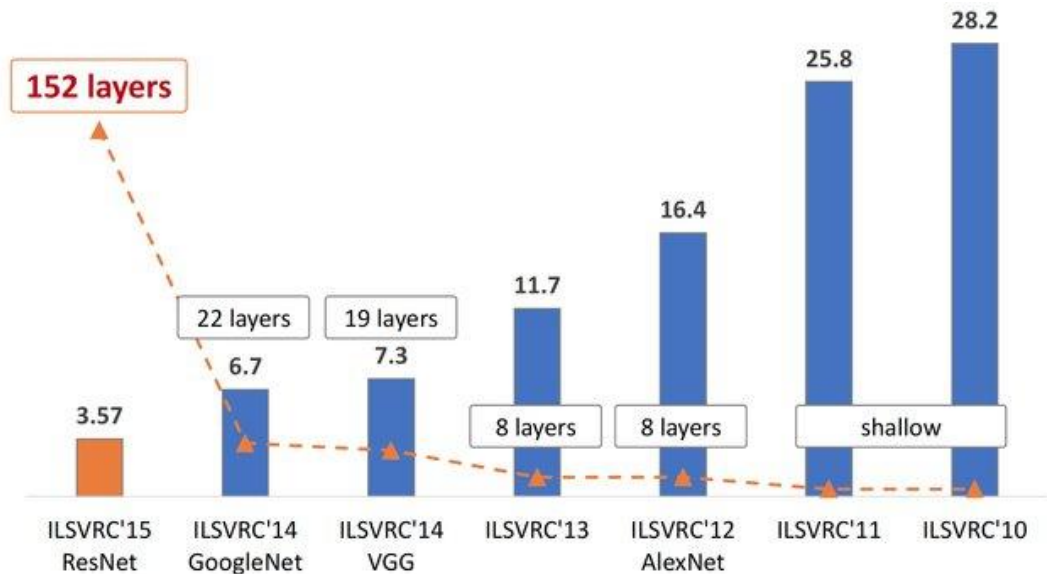
# ImageNet



Everything starts with data!

- The ImageNet dataset contains 14,197,122 annotated images.
- Total number of non-empty WordNet synsets: 21,841.
- Since 2010 the dataset is used in the ImageNet Large Scale Visual Recognition Challenge (ILSVRC), a benchmark in image classification and object detection.

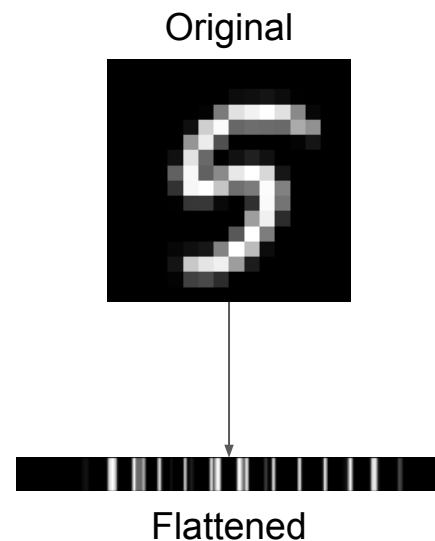
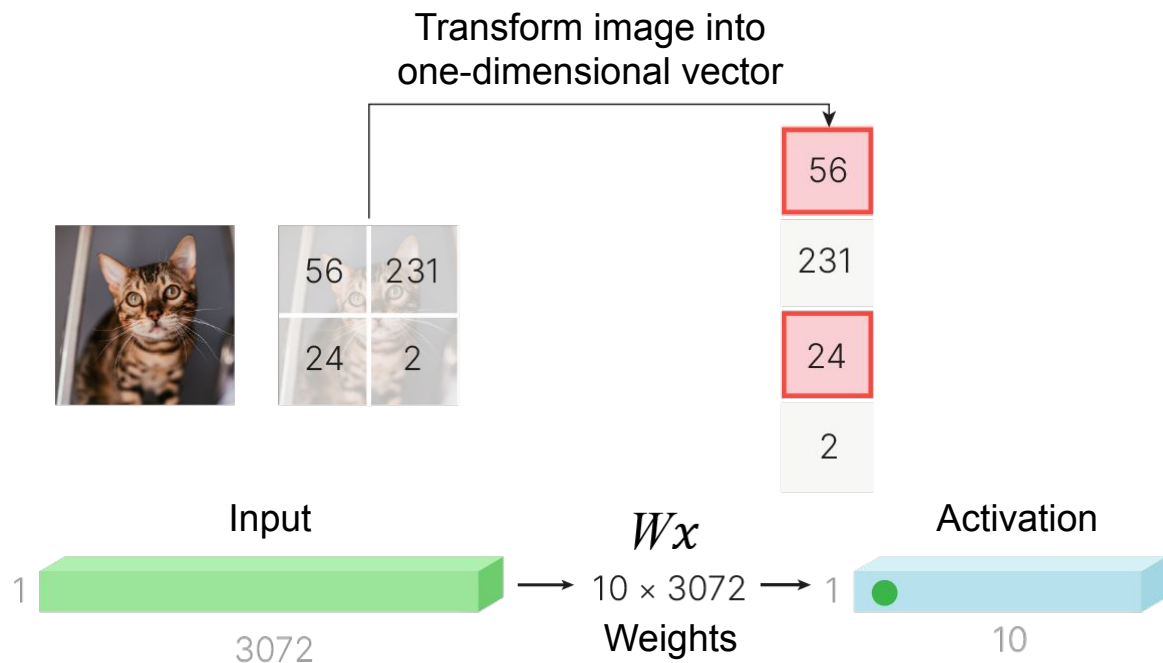
# ImageNet Results



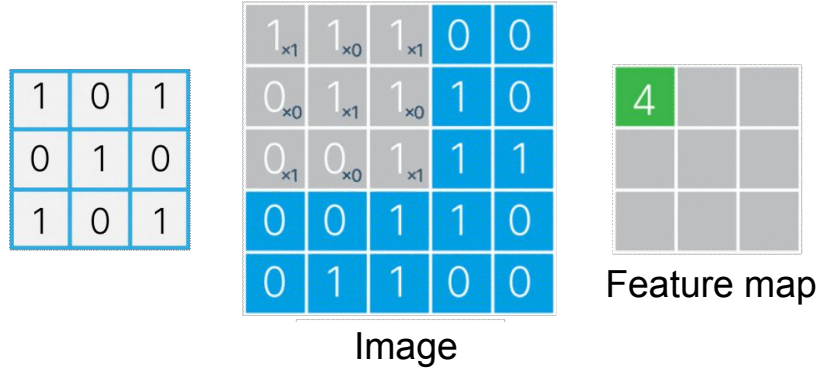
Why we can't use MLP for images?

- Too many parameters
- Fixed dimension of images
- Features will be too correlated

# Inherent problem of MLP



# Sliding window (filter)



Original



Gaussian



Sobel



# Motivation for convolution

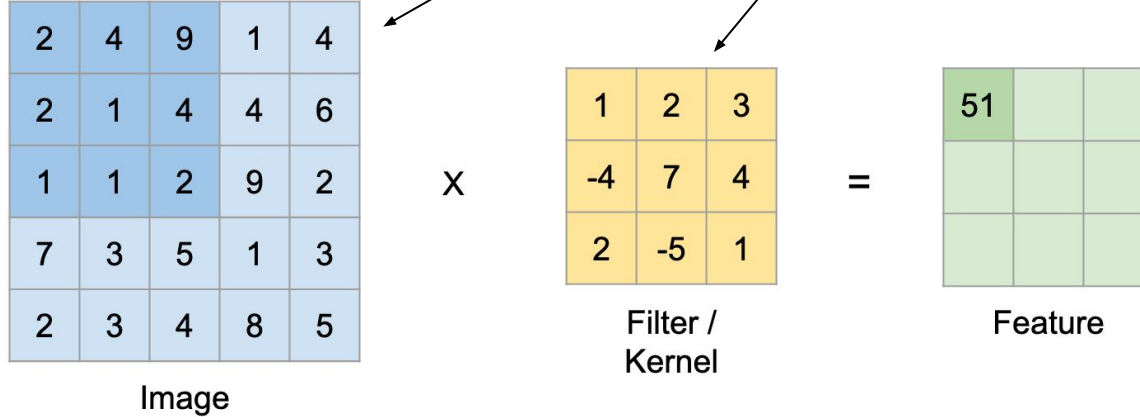
- To perform identification, we are seeking for noticeable parts of an object
- However, these parts appear at different locations
- Want to perform a feature search over the whole picture
- **Idea:** Let the neural network choose the filters by itself



# Convolution

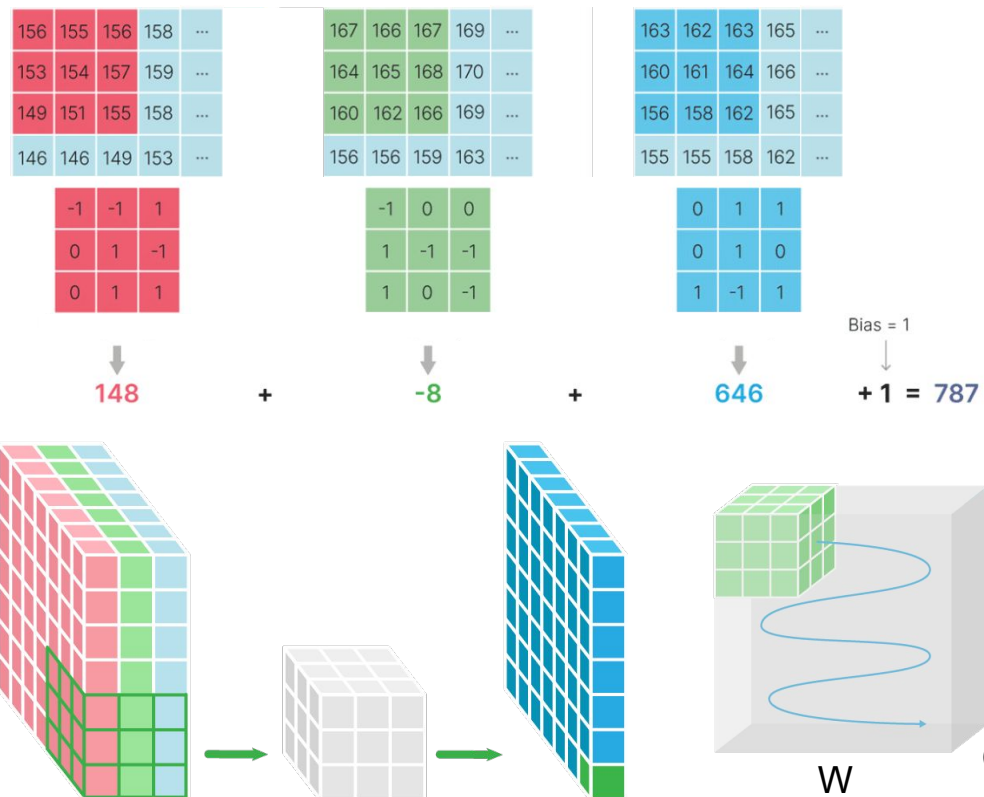
Mathematically, 2D convolution\* can be written in the following form

$$Y(i, j) = \sum_{u, v} X(i + u, j + v) K(u, v)$$



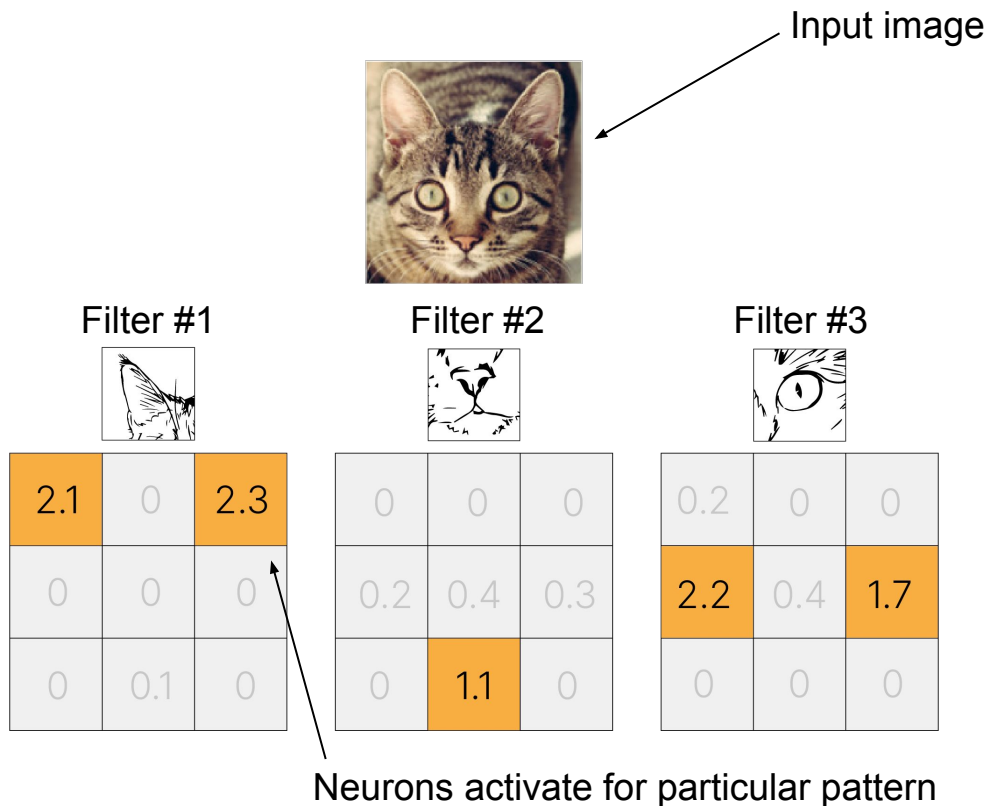
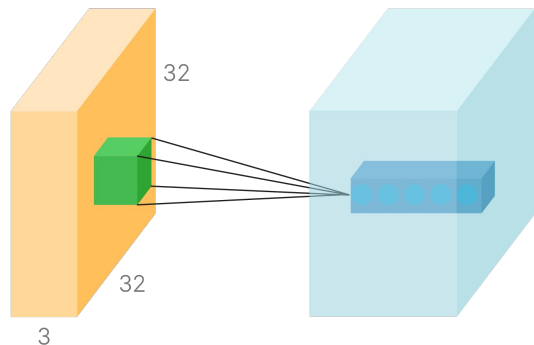
\*How to calculate the gradient? See [this post](#) on Pavithra Solai Blog

## Color images / multiple input channels



# Multiple output channels

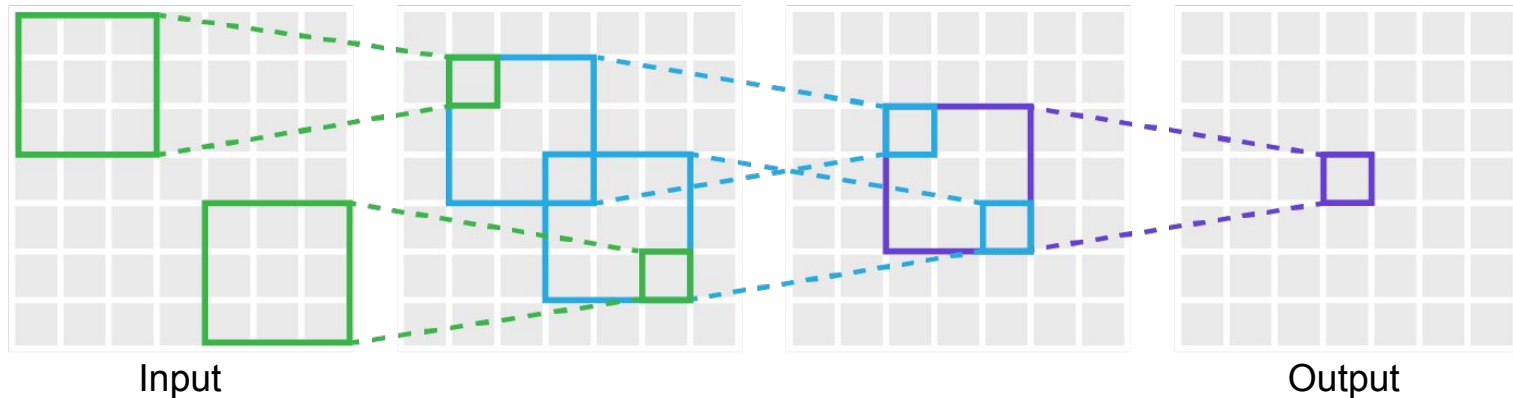
- In fully-connected neural networks, each neuron activates for particular pattern (red car, blue chair, etc.)
- We want CNN neurons to do the same, e.g. ear, nose, eye
- Thus, for **each pattern** we should have a neuron, i.e. a **filter**



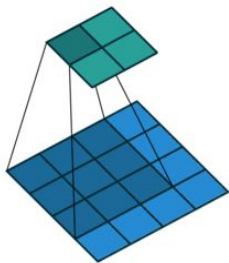


# Receptive field

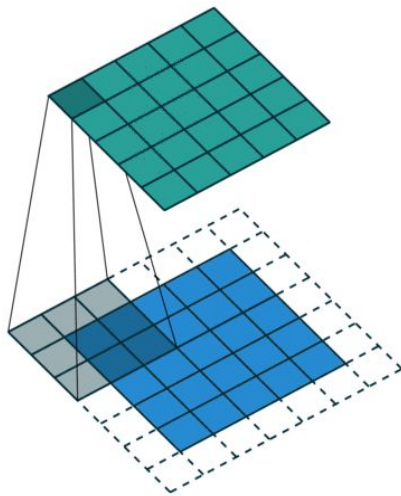
- Usually complex features consist of less complex features: e.g. a cat has a face, the face has ears, the ears have edges
- CNN layers, close to the:
  - Input → detect less complex features
  - Output → detect more complex features, combining previous ones
- Receptive field is the input image region that a particular neuron is “looking at”



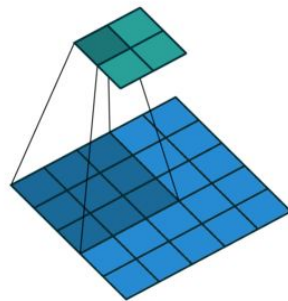
# How to save dimension / increase receptive field?



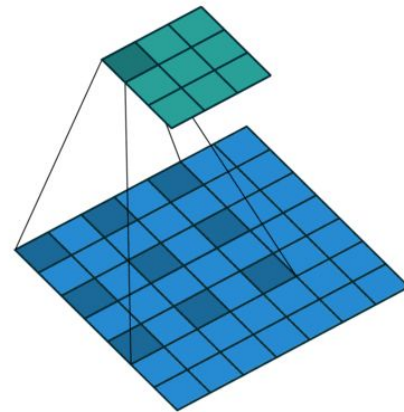
No padding,  
no strides



Padding 1,  
no strides

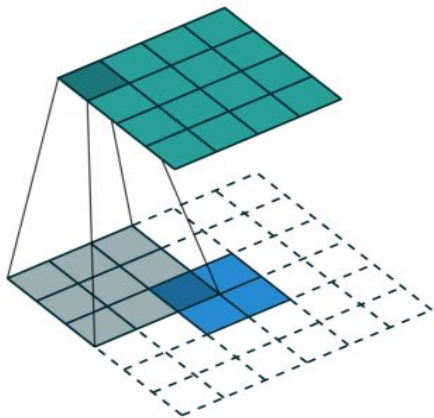


No padding,  
stride = 2

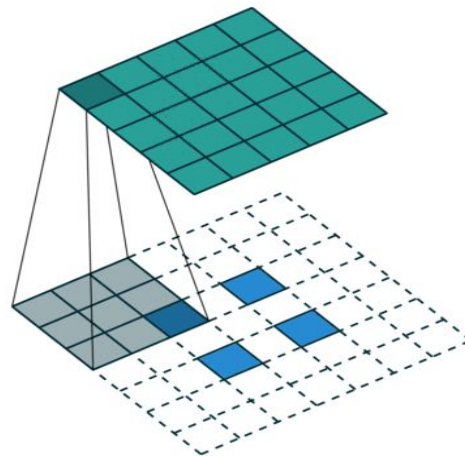


No padding,  
no stride,  
dilation = 2

# Upsampling with convolutional layer

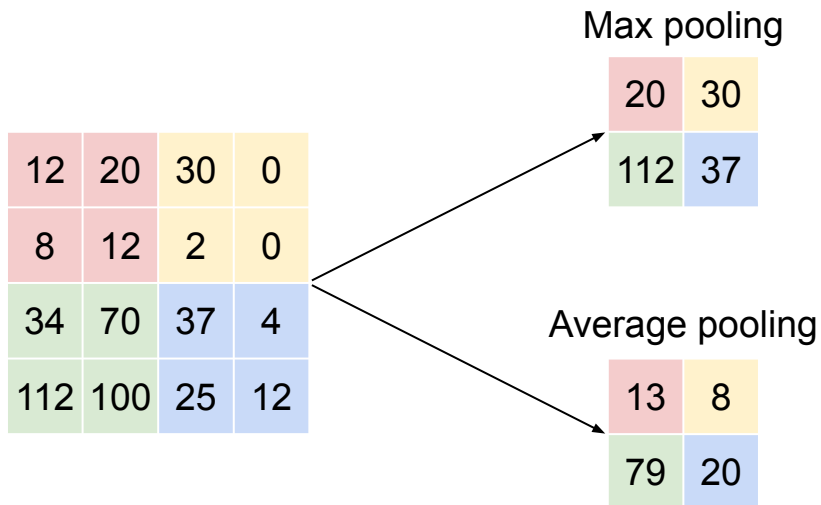


Transposed convolution



Transposed convolution  
stride = 2

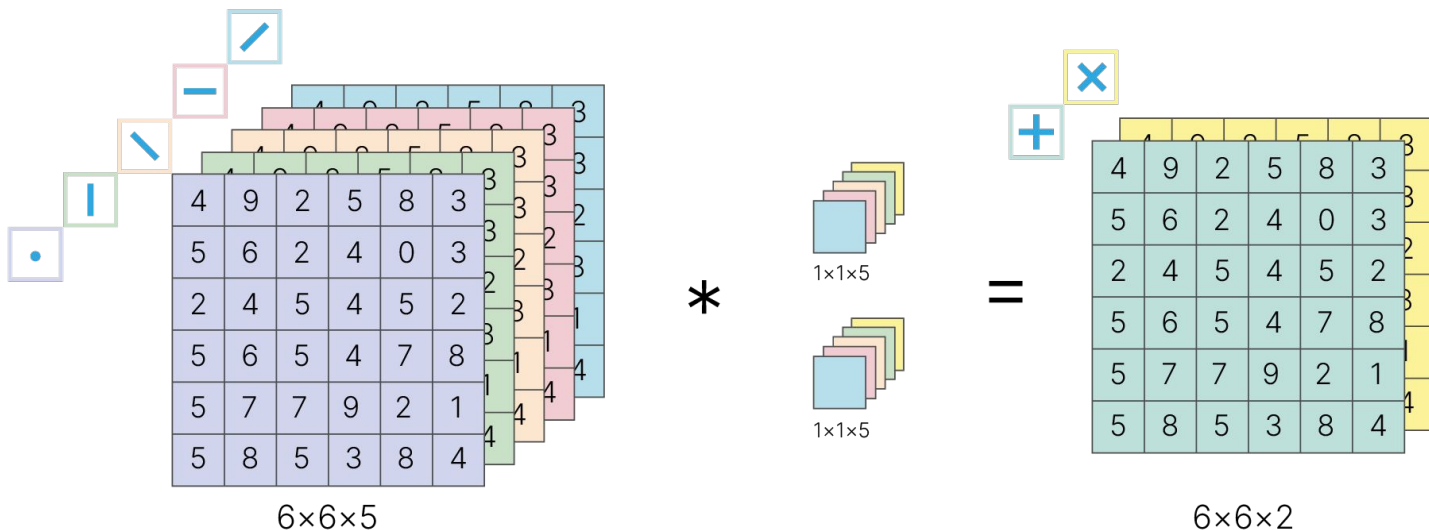
# Pooling



- Usual motivation: adding invariance to small shifts
- Several max-poolings can accumulate invariance to stronger shifts
- Pooling is applied for each channel, so the total number of channels remains unchanged

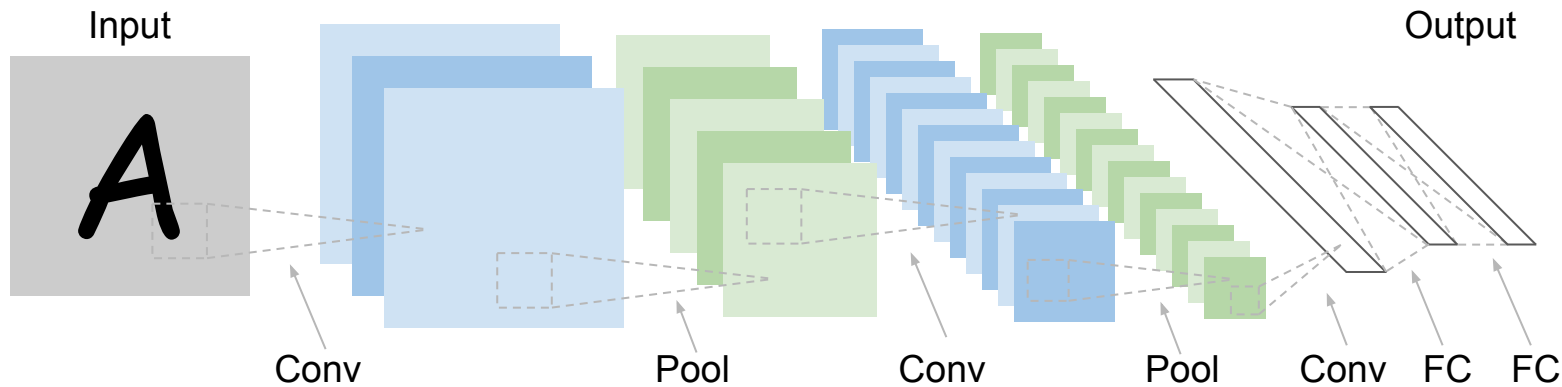
# Convolution 1x1

- Reducing the number of channels leads to a generalization of features
- It is often obtained through the 1x1 convolutional layer
- Example: RGB  $\rightarrow$  Grayscale image, i.e.  $\text{Brightness} = \frac{1}{3}R + \frac{1}{3}G + \frac{1}{3}B$



# LeNet: example of CNN

- In the 1989, Yann LeCun et al. introduced the LeNet-5 neural network
- Character recognition: handwriting and machine printing
- Seven layers: 3 convolutional, 2 pooling, 2 fully-connected
- Activation function: tanh



# Recap

- Vanishing and exploding gradient
- Weights initialization
- Internal Covariate Shift
- Normalization
  - Batch Norm
  - Layer Norm
  - Instance Norm
  - Group Norm
- Convolutional NN



**TIME FOR A BREAK**