

Deep Learning

Lecture 9.2

reused a parts of a good (theory) course at HSE University (http://wiki.cs.hse.ru/Reinforcement_learning_2022_2023)

In previous lecture



- Markov Decision Processes
- Value Iteration -> Least-Squares Value Iteration
- Exploration-Exploitation Tradeoff
- Experience Replay (Replay Buffer)
- Deep Q-Network (DQN)
- Atari & Procgen Benchmarks

Recap: RL and Markov Decision Process

Markov Decision Process is a 5-tuple (S, A, P, R, γ)

- S - state space;
- A - action space;
- $P(s' | s, a)$ - transition probability kernel;
- $R(s, a)$ - reward distribution (with a mean reward $r(s, a)$);
- γ - discounting factor

Policy π : rule to choose next action given current state



MDP Example: Chain

Recap: Value function and Q-function

Goal of the agent: find a policy π that maximizes the expected sum of rewards:

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right] \quad Q^\pi(s, a) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a \right],$$

$$r_t \sim R(\cdot | s_t, a_t), s_{t+1} \sim P(\cdot | s_t, a_t), a_t \sim \pi(\cdot | s_t).$$

A policy that attains maximum for each states is called *optimal*.

RL Objective

- Goal of RL – find good policy, so let us parameterize policies!

$$\pi_{\theta}(a|s) = \frac{\exp(f_{\theta}(s, a))}{\sum_{a' \in \mathcal{A}} \exp(f_{\theta}(s, a'))}$$

- Let s_0 be an initial state. Then our objective is $J(\theta) = V^{\pi_{\theta}}(s_0)$

Overall, we recast RL problem as **optimization problem**

$$\max_{\theta} J(\theta)$$

Q: Why this problem is difficult?

Policy Optimization – first difficulties

- Recall the definition of value:

$$V^\pi(s) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right],$$

- Policy is **hidden** inside the expectation in a very non-direct way
- How to compute gradients with respect to the policy to perform gradient ascent (for example)?

$$\theta_{t+1} = \theta_t + \alpha_t \nabla J(\theta_t)$$

Policy Gradient Theorem

Theorem 15 (Policy Gradient Theorem). *Let $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ be discounted MDP. Let $B: \mathcal{S} \rightarrow \mathbb{R}$ be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter θ is equal to*

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

- We will discuss what is baseline and how to choose it later.
Now just think that $B = 0$.

Proof of Policy Gradient Theorem, pt.1

By Bellman equations $V^{\pi_\theta}(s_0) = \sum_{a_0 \in \mathcal{A}} \pi_\theta(a_0|s_0) Q^{\pi_\theta}(s, a),$

By chain rule:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \sum_{a_0 \in \mathcal{A}} [\nabla_\theta \pi_\theta(a_0|s_0) Q^{\pi_\theta}(s_0, a_0) + \pi_\theta(a_0|s_0) \cdot \nabla Q^{\pi_\theta}(s_0, a_0)].$$

“Log-derivative trick”: $\nabla_\theta \log \pi_\theta(a|s) = \frac{\nabla_\theta \pi_\theta(a|s)}{\pi_\theta(a|s)},$

Overall:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{a_0 \sim \pi_\theta(\cdot|s_0)} [\nabla_\theta \log \pi_\theta(a_0|s_0) \cdot Q^{\pi_\theta}(s_0, a_0) + \nabla Q^{\pi_\theta}(s_0, a_0)].$$

Proof of Policy Gradient Theorem, pt.2

Bellman equations:

$$Q^\pi(s, a) = r(s, a) + \gamma P V^\pi(s, a),$$

$$V^\pi(s) = \sum_{a \in \mathcal{A}} Q^\pi(s, a) \pi(a, s)$$

As a result:

$$\nabla_\theta Q^{\pi_\theta}(s_0, a_0) = \nabla_\theta [r(a_0, s_0) + \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [V^{\pi_\theta}(s_1)]] = \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [\nabla_\theta V^{\pi_\theta}(s_1)],$$

Plug-in in derivative for value:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{a_0 \sim \pi_\theta(\cdot | s_0), s_1 \sim P(\cdot | s_0, a_0)} [\nabla_\theta \log \pi_\theta(a_0 | s_0) \cdot Q^{\pi_\theta}(s_0, a_0) + \gamma \nabla_\theta V^{\pi_\theta}(s_1)].$$

Rolling-out we obtain

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot Q^{\pi_\theta}(s_t, a_t) \right],$$

Proof of Policy Gradient Theorem, pt.3

To finish the proof, we have to show $\mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t) \cdot B(s_t) \right] = 0.$

$$\begin{aligned} \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t) \cdot B(s_t) \right] &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0), a \sim \pi_\theta(\cdot|s_t)} [\nabla_\theta \log \pi_\theta(a|s_t) \cdot B(s_t)] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[\sum_{a \in \mathcal{A}} \pi_\theta(a|s_t) \nabla_\theta \log \pi_\theta(a|s_t) \cdot B(s_t) \right] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[\sum_{a \in \mathcal{A}} \nabla_\theta \pi_\theta(a|s_t) \cdot B(s_t) \right] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[\nabla_\theta \left(\underbrace{\sum_{a \in \mathcal{A}} \pi_\theta(a|s_t)}_1 \right) \cdot B(s_t) \right] = 0. \end{aligned}$$

Policy Gradient Theorem

Theorem 15 (Policy Gradient Theorem). *Let $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ be discounted MDP. Let $B: \mathcal{S} \rightarrow \mathbb{R}$ be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter θ is equal to*

equal to

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

unknown! (pointing to π_{θ})

unknown! (pointing to $Q^{\pi_{\theta}}$)

- **Q:** How to apply it in the real life?

Policy Gradient Estimation

- Assume that we used a current policy to obtain trajectory

$$(s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots).$$

- Estimate Q-value: $Q^{\pi_\theta}(s_t, a_t) \approx G_t = \sum_{k=t}^{\infty} \gamma^{k-t} r_k.$

- Estimate outer expectation $\hat{\nabla} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot G_t.$

Lemma 5. $\hat{\nabla} J(\theta)$ is an unbiased estimate of $\nabla J(\theta)$.

REINFORCE

Algorithm that uses the following gradient estimate to perform SGD is called REINFORCE

Algorithm 6 REINFORCE

Input: MDP $M = (\mathcal{S}, \mathcal{A}, P, R, H)$;

Initialize: θ_0 ;

for $k = 0, 1, \dots$, **do**

 Play policy π_{θ_k} and receive a trajectory $s_0^k, a_0^k, r_0^k, \dots, s_T^k, a_T^k, r_T^k$.

 Compute estimates of Q-function $G_t = \sum_{i=t}^T \gamma^{i-t} r_i$ for all $t = 0, \dots, T$;

 Compute estimate of policy gradient $\hat{\nabla} J(\theta_k) = \sum_{t=0}^T \gamma^t \nabla_{\theta} \pi_{\theta}(a_t | s_t) \cdot G_t$;

 Perform a stochastic gradient step $\theta_{k+1} = \theta_k + \alpha_k \hat{\nabla} J(\theta_k)$.

end for

Remark. Algorithm can utilize only trajectories that were collected during the training, such algorithms are called *on-policy*.

Algorithms that can utilize data obtained by other policy are called *off-policy*.

Problems of REINFORCE

- Inefficient sample utilization;
 - Requires fast simulator!
- Large variance that leads to slow convergence;
 - Can be handled somehow by working with several parallel environments and collecting trajectories.

Variance Reduction: Actor-Critic Algorithm

Idea 1. Let's estimate Q-value smarter, we have Bellman equations for it!

$$Q^\pi(s, a) = r(s, a) + \gamma P V^\pi(s, a),$$

$$V^\pi(s) = \sum_{a \in \mathcal{A}} Q^\pi(s, a) \pi(a, s)$$

We can do it in DQN-fashion! $Q_\psi \approx Q^{\pi_\theta},$

$$\hat{\nabla}_{\text{AC}} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot Q_{\psi}(s_t, a_t)$$

Remember this guy? Connection to policy iteration

Algorithm 1 Policy Iteration for discounted MDPs

Input: MDP $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$, the immediate reward function r , iterations budget T

Initialize: π^0 as some set of policies;

for $t \in [T]$ **do**

 Compute Q^{π^t} by solving Bellman equations (see Theorem 4);


 Find π^{t+1} as a greedy policy w.r.t. Q^{π^t} .

end for

Output: estimate of optimal policy π^T .

Policy Evaluation
steps

Policy Improvement steps



Variance Reduction: baseline selection

Idea 2. Use baseline!

Theorem 15 (Policy Gradient Theorem). *Let $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ be discounted MDP. Let $B: \mathcal{S} \rightarrow \mathbb{R}$ be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter θ is equal to*

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

Natural choice: $B(s)$ is value function!

(this choice is not unique and is not optimal!)

Advantage Actor Critic (A2C)

Remark 2. The baseline is needed for variance reduction purposes. The most common choice is $B(s) = V^{\pi_\theta}(s)$ that leads to the following view on policy gradient theorem

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot A^{\pi_\theta}(s_t, a_t) \right],$$

where $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$ is *advantage* function.

Q: How to estimate advantage function?

Advantage estimation

First, let us provide unbiased estimate:

$$\begin{aligned} A^{\pi_{\theta}}(s_t, a_t) &= Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t) \\ &= \mathbb{E}_{r_t \sim R(s_t, a_t) s'_t \sim P(s_t, a_t)} [r_t + \gamma V^{\pi_{\theta}}(s'_t)] - V^{\pi_{\theta}}(s_t) \\ &\approx r_t + \gamma V^{\pi_{\theta}}(s_{t+1}) - V^{\pi_{\theta}}(s_t) \end{aligned}$$

In this estimation we need a network only for V-function, that is usually much easier to learn! We can do it through Bellman equations and optimizing TD loss

$$\mathcal{L}_{\text{critic}}(\psi) = \sum_{t=1}^T \left(V_{\psi}(s_t) - r_t - \gamma V_{\tilde{\psi}}(s_{t+1}) \right)^2$$

target network, usually
just stopgrad here

A2C as Policy Iteration

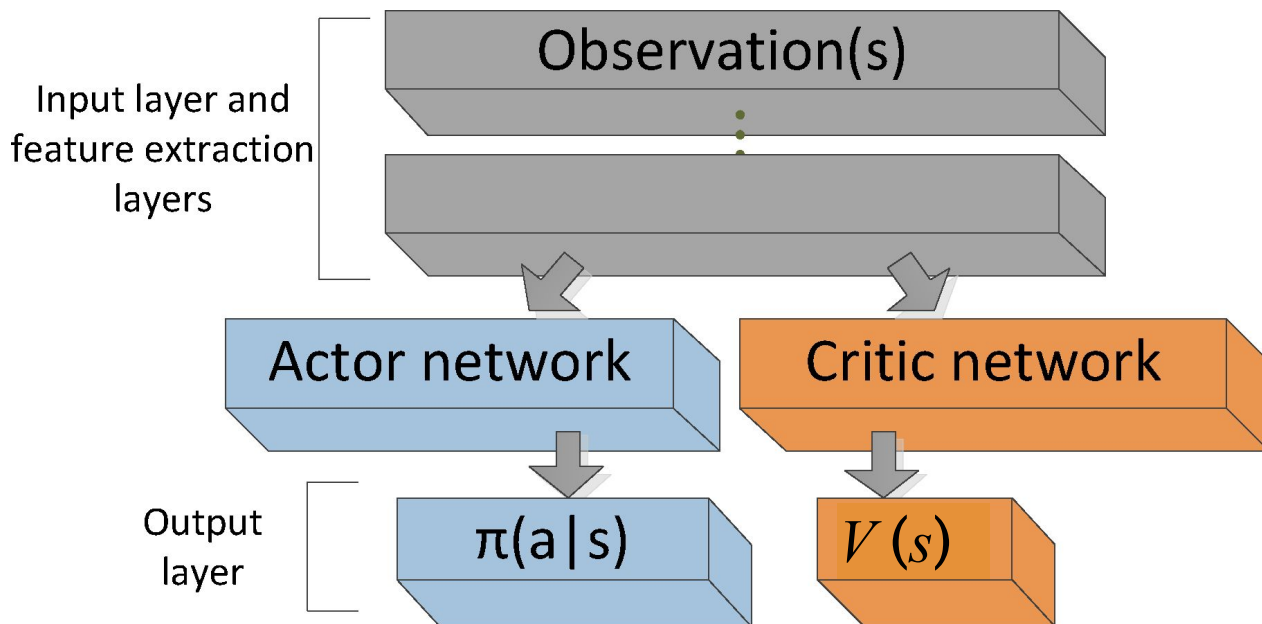
Policy Improvement: step by

$$\hat{\nabla}_{\text{A2C}} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (r_t + \gamma V_{\psi}(s_{t+1}) - V_{\psi}(s_t))$$

Policy evaluation: gradient step on the following loss

$$\mathcal{L}_{\text{critic}}(\psi) = \sum_{t=1}^T \left(V_{\psi}(s_t) - r_t - \gamma V_{\tilde{\psi}}(s_{t+1}) \right)^2$$

Architecture Details



source: AI Masters RL course (<https://ozonmasters.ru/reinforcementlearning>)

Exploration-Exploitation Trade-off

For DQN we need exploration to satisfy constraints on the replay buffer generation distribution.

Q: Do we need additional exploration for policy gradient methods?

We just reformulate RL problem as an optimization problem and perform SGD!

$$\max_{\theta} J(\theta)$$


Exploration for Policy-Gradient methods

A: We still need exploration since we can stuck in a local optimum!

One common way: add negative entropy to the loss function for actor network to maximize it!

$$\mathcal{H}(\pi_{\theta}(s_t)) = \sum_{a \in \mathcal{A}} \pi_{\theta}(a|s_t) \log\left(\frac{1}{\pi_{\theta}(a|s_t)}\right)$$

- Encourages exploration;
- Introduces bias, so a coefficient should be small!



Advanced topics on Policy Gradient methods

- Generalized Advantage Estimation (GAE)
- Introduction of a small delay into on-policy generation
 - Proximal Policy Optimization (PPO);
 - Trust-Region Policy Optimization (TRPO);
 - Asynchronous Advantage Actor Critic (A3C) and Impala;
 - Mirror Descent Policy Optimization (MDPO);
- Different types of baselines:
 - Hindsight Credit Assignment;
 - Action-dependent baselines;
- IMPLEMENTATION MATTERS IN DEEP POLICY GRADIENTS: A CASE STUDY ON PPO AND TRPO

<https://arxiv.org/abs/2005.12729>



Recap for RL

- What is RL?
- Markov Decision Process;
- RL problems
- V-function, Q-function;
- Value Iteration, Policy iteration;
- Least-Squared Value Iteration;
- Exploration-Exploitation trade-off;
- Experience Replay (Replay Buffer);
- Deep Q-Network (DQN);
- Policy Gradient Theorem, log-derivative trick;
- REINFORCE;
- Actor-Critic algorithm and A2C;