

# Deep Learning

## Lecture 10

reused a parts of a good (theory) course at HSE University ([http://wiki.cs.hse.ru/Reinforcement\\_learning\\_2022\\_2023](http://wiki.cs.hse.ru/Reinforcement_learning_2022_2023))

Lecturer: Daniil Tiapkin (<https://d-tiapkin.github.io/>)

# In previous lecture



- Markov Decision Processes
- Value Iteration -> Least-Squares Value Iteration
- Exploration-Exploitation Tradeoff
- Experience Replay (Replay Buffer)
- Deep Q-Network (DQN)
- Atari & Procgen Benchmarks

# Recap: RL and Markov Decision Process

**Markov Decision Process** is a 5-tuple  $(S, A, P, R, \gamma)$

- $S$  - state space;
- $A$  - action space;
- $P(s' | s, a)$  - transition probability kernel;
- $R(s, a)$  - reward distribution (with a mean reward  $r(s, a)$ );
- $\gamma$  - discounting factor

**Policy  $\pi$ :** rule to choose next action given current state



MDP Example: Chain

# Recap: Value function and Q-function

**Goal of the agent:** find a policy  $\pi$  that maximizes the expected sum of rewards:

$$V^{\pi}(s) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right] \quad Q^{\pi}(s, a) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s, a_0 = a \right],$$

$$r_t \sim R(\cdot | s_t, a_t), s_{t+1} \sim P(\cdot | s_t, a_t), a_t \sim \pi(\cdot | s_t).$$

A policy that attains maximum for each states is called *optimal*.

# RL Objective

- Goal of RL – find good policy, so let us parameterize policies!

$$\pi_{\theta}(a|s) = \frac{\exp(f_{\theta}(s, a))}{\sum_{a' \in \mathcal{A}} \exp(f_{\theta}(s, a'))}$$

- Let  $s_0$  be an initial state. Then our objective is  $J(\theta) = V^{\pi_{\theta}}(s_0)$

Overall, we recast RL problem as **optimization problem**

$$\max_{\theta} J(\theta)$$

**Q:** Why this problem is difficult?

# Policy Optimization – first difficulties

- Recall the definition of value:

$$V^\pi(s) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t r_t \mid s_0 = s \right],$$

- Policy is **hidden** inside the expectation in a very non-direct way
- How to compute gradients with respect to the policy to perform gradient ascent (for example)?

$$\theta_{t+1} = \theta_t + \alpha_t \nabla J(\theta_t)$$

# Policy Gradient Theorem

**Theorem 15** (Policy Gradient Theorem). *Let  $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$  be discounted MDP. Let  $B: \mathcal{S} \rightarrow \mathbb{R}$  be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter  $\theta$  is equal to*

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

- We will discuss what is baseline and how to choose it later.  
Now just think that  $B = 0$ .

# Proof of Policy Gradient Theorem, pt.1

By Bellman equations  $V^{\pi_\theta}(s_0) = \sum_{a_0 \in \mathcal{A}} \pi_\theta(a_0|s_0) Q^{\pi_\theta}(s, a),$

By chain rule:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \sum_{a_0 \in \mathcal{A}} [\nabla_\theta \pi_\theta(a_0|s_0) Q^{\pi_\theta}(s_0, a_0) + \pi_\theta(a_0|s_0) \cdot \nabla Q^{\pi_\theta}(s_0, a_0)].$$

“Log-derivative trick”:  $\nabla_\theta \log \pi_\theta(a|s) = \frac{\nabla_\theta \pi_\theta(a|s)}{\pi_\theta(a|s)},$

Overall:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{a_0 \sim \pi_\theta(\cdot|s_0)} [\nabla_\theta \log \pi_\theta(a_0|s_0) \cdot Q^{\pi_\theta}(s_0, a_0) + \nabla Q^{\pi_\theta}(s_0, a_0)].$$



# Proof of Policy Gradient Theorem, pt.2

Bellman equations:

$$Q^\pi(s, a) = r(s, a) + \gamma P V^\pi(s, a),$$

$$V^\pi(s) = \sum_{a \in \mathcal{A}} Q^\pi(s, a) \pi(a, s)$$

As a result:

$$\nabla_\theta Q^{\pi_\theta}(s_0, a_0) = \nabla_\theta [r(a_0, s_0) + \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [V^{\pi_\theta}(s_1)]] = \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [\nabla_\theta V^{\pi_\theta}(s_1)],$$

Plug-in in derivative for value:

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{a_0 \sim \pi_\theta(\cdot | s_0), s_1 \sim P(\cdot | s_0, a_0)} [\nabla_\theta \log \pi_\theta(a_0 | s_0) \cdot Q^{\pi_\theta}(s_0, a_0) + \gamma \nabla_\theta V^{\pi_\theta}(s_1)].$$

Rolling-out we obtain

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot Q^{\pi_\theta}(s_t, a_t) \right],$$

## Proof of Policy Gradient Theorem, pt.3

To finish the proof, we have to show  $\mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t) \cdot B(s_t) \right] = 0.$

$$\begin{aligned} \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t|s_t) \cdot B(s_t) \right] &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0), a \sim \pi_\theta(\cdot|s_t)} [\nabla_\theta \log \pi_\theta(a|s_t) \cdot B(s_t)] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[ \sum_{a \in \mathcal{A}} \pi_\theta(a|s_t) \nabla_\theta \log \pi_\theta(a|s_t) \cdot B(s_t) \right] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[ \sum_{a \in \mathcal{A}} \nabla_\theta \pi_\theta(a|s_t) \cdot B(s_t) \right] \\ &= \sum_{t=0}^{\infty} \gamma^t \mathbb{E}_{s_t \sim (P^{\pi_\theta})^t(\cdot|s_0)} \left[ \nabla_\theta \left( \underbrace{\sum_{a \in \mathcal{A}} \pi_\theta(a|s_t)}_1 \right) \cdot B(s_t) \right] = 0. \end{aligned}$$

# Policy Gradient Theorem

**Theorem 15** (Policy Gradient Theorem). *Let  $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$  be discounted MDP. Let  $B: \mathcal{S} \rightarrow \mathbb{R}$  be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter  $\theta$  is equal to*

equal to

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

unknown! (pointing to  $\pi_{\theta}$ )

unknown! (pointing to  $Q^{\pi_{\theta}}$ )

- **Q:** How to apply it in the real life?

# Policy Gradient Estimation

- Assume that we used a current policy to obtain trajectory

$$(s_0, a_0, r_0, \dots, s_t, a_t, r_t, \dots).$$

- Estimate Q-value:  $Q^{\pi_\theta}(s_t, a_t) \approx G_t = \sum_{k=t}^{\infty} \gamma^{k-t} r_k.$

- Estimate outer expectation  $\hat{\nabla} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot G_t.$

**Lemma 5.**  $\hat{\nabla} J(\theta)$  is an unbiased estimate of  $\nabla J(\theta)$ .

# REINFORCE

Algorithm that uses the following gradient estimate to perform SGD is called REINFORCE

---

**Algorithm 6** REINFORCE

---

**Input:** MDP  $M = (\mathcal{S}, \mathcal{A}, P, R, H)$ ;

**Initialize:**  $\theta_0$ ;

**for**  $k = 0, 1, \dots$ , **do**

    Play policy  $\pi_{\theta_k}$  and receive a trajectory  $s_0^k, a_0^k, r_0^k, \dots, s_T^k, a_T^k, r_T^k$ .

    Compute estimates of Q-function  $G_t = \sum_{i=t}^T \gamma^{i-t} r_i$  for all  $t = 0, \dots, T$ ;

    Compute estimate of policy gradient  $\hat{\nabla} J(\theta_k) = \sum_{t=0}^T \gamma^t \nabla_{\theta} \pi_{\theta}(a_t | s_t) \cdot G_t$ ;

    Perform a stochastic gradient step  $\theta_{k+1} = \theta_k + \alpha_k \hat{\nabla} J(\theta_k)$ .

**end for**

---

**Remark.** Algorithm can utilize only trajectories that were collected during the training, such algorithms are called *on-policy*.

Algorithms that can utilize data obtained by other policy are called *off-policy*.

# Problems of REINFORCE

- Inefficient sample utilization;
  - Requires fast simulator!
- Large variance that leads to slow convergence;
  - Can be handled somehow by working with several parallel environments and collecting trajectories.

# Variance Reduction: Actor-Critic Algorithm

**Idea 1.** Let's estimate Q-value smarter, we have Bellman equations for it!

$$Q^\pi(s, a) = r(s, a) + \gamma P V^\pi(s, a),$$

$$V^\pi(s) = \sum_{a \in \mathcal{A}} Q^\pi(s, a) \pi(a, s)$$

We can do it in DQN-fashion!  $Q_\psi \approx Q^{\pi_\theta},$

$$\hat{\nabla}_{\text{AC}} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot Q_{\psi}(s_t, a_t)$$

# Remember this guy? Connection to policy iteration

---

**Algorithm 1** Policy Iteration for discounted MDPs

---

**Input:** MDP  $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$ , the immediate reward function  $r$ , iterations budget  $T$

**Initialize:**  $\pi^0$  as some set of policies;

**for**  $t \in [T]$  **do**

    Compute  $Q^{\pi^t}$  by solving Bellman equations (see Theorem 4);

    Find  $\pi^{t+1}$  as a greedy policy w.r.t.  $Q^{\pi^t}$ .


**end for**

**Output:** estimate of optimal policy  $\pi^T$ .

---

Policy Evaluation  
steps

Policy Improvement steps





# Variance Reduction: baseline selection

## Idea 2. Use baseline!

**Theorem 15** (Policy Gradient Theorem). *Let  $M = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$  be discounted MDP. Let  $B: \mathcal{S} \rightarrow \mathbb{R}$  be any bounded function on states (so called baseline). Then the gradient of value function with respect to parameter  $\theta$  is equal to*

$$\nabla_{\theta} V^{\pi_{\theta}}(s_0) = \mathbb{E}_{\pi_{\theta}} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (Q^{\pi_{\theta}}(s_t, a_t) - B(s_t)) \right].$$

**Natural choice:**  $B(s)$  is value function!

(this choice is not unique and is not optimal!)

# Advantage Actor Critic (A2C)

**Remark 2.** The baseline is needed for variance reduction purposes. The most common choice is  $B(s) = V^{\pi_\theta}(s)$  that leads to the following view on policy gradient theorem

$$\nabla_\theta V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^{\infty} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot A^{\pi_\theta}(s_t, a_t) \right],$$

where  $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$  is *advantage* function.

**Q:** How to estimate advantage function?

# Advantage estimation

First, let us provide unbiased estimate:

$$\begin{aligned} A^{\pi_{\theta}}(s_t, a_t) &= Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t) \\ &= \mathbb{E}_{r_t \sim R(s_t, a_t) s'_t \sim P(s_t, a_t)} [r_t + \gamma V^{\pi_{\theta}}(s'_t)] - V^{\pi_{\theta}}(s_t) \\ &\approx r_t + \gamma V^{\pi_{\theta}}(s_{t+1}) - V^{\pi_{\theta}}(s_t) \end{aligned}$$

In this estimation we need a network only for V-function, that is usually much easier to learn! We can do it through Bellman equations and optimizing TD loss

$$\mathcal{L}_{\text{critic}}(\psi) = \sum_{t=1}^T \left( V_{\psi}(s_t) - r_t - \gamma V_{\tilde{\psi}}(s_{t+1}) \right)^2$$

target network, usually  
just stopgrad here

# A2C as Policy Iteration

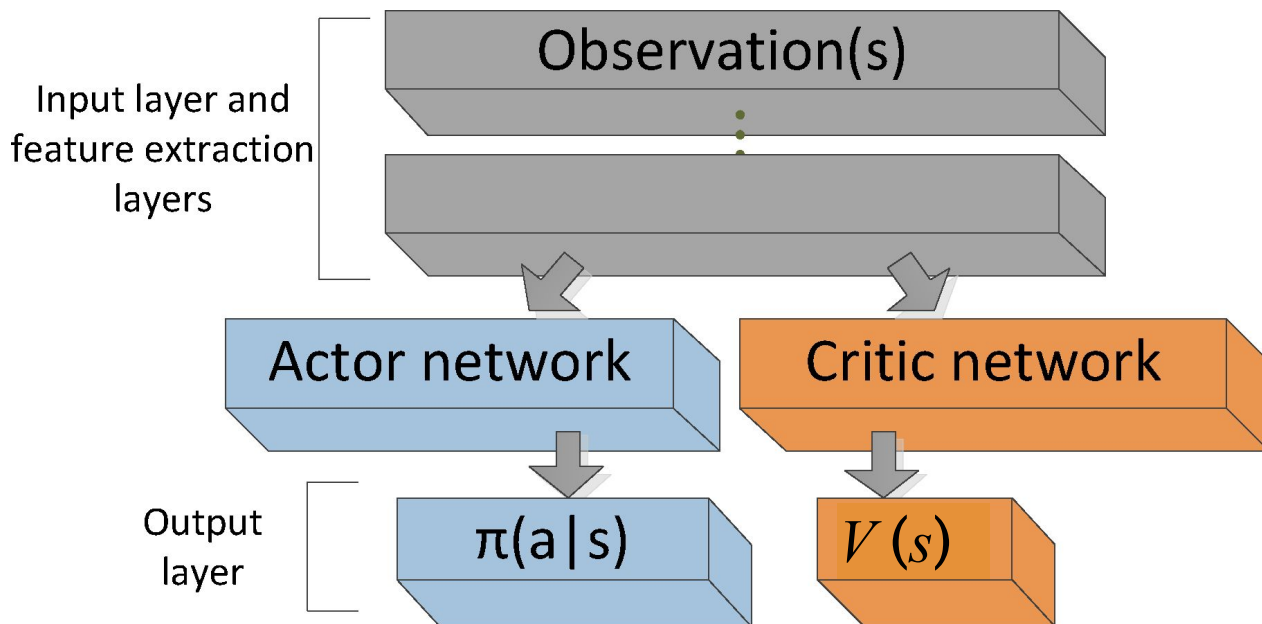
**Policy Improvement:** step by

$$\hat{\nabla}_{\text{A2C}} J(\theta) = \sum_{t=0}^{\infty} \gamma^t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot (r_t + \gamma V_{\psi}(s_{t+1}) - V_{\psi}(s_t))$$

**Policy evaluation:** gradient step on the following loss

$$\mathcal{L}_{\text{critic}}(\psi) = \sum_{t=1}^T \left( V_{\psi}(s_t) - r_t - \gamma V_{\tilde{\psi}}(s_{t+1}) \right)^2$$

# Architecture Details



source: AI Masters RL course (<https://ozonmasters.ru/reinforcementlearning>)

# Exploration-Exploitation Trade-off

For DQN we need exploration to satisfy constraints on the replay buffer generation distribution.

**Q:** Do we need additional exploration for policy gradient methods?

We just reformulate RL problem as an optimization problem and perform SGD!

$$\max_{\theta} J(\theta)$$


# Exploration for Policy-Gradient methods

**A:** We still need exploration since we can stuck in a local optimum!

One common way: add negative entropy to the loss function for actor network to maximize it!

$$\mathcal{H}(\pi_{\theta}(s_t)) = \sum_{a \in \mathcal{A}} \pi_{\theta}(a|s_t) \log\left(\frac{1}{\pi_{\theta}(a|s_t)}\right)$$

- Encourages exploration;
- Introduces bias, so a coefficient should be small!



# Advanced topics on Policy Gradient methods

- Generalized Advantage Estimation (GAE)
- Introduction of a small delay into on-policy generation
  - Proximal Policy Optimization (PPO);
  - Trust-Region Policy Optimization (TRPO);
  - Asynchronous Advantage Actor Critic (A3C) and Impala;
  - Mirror Descent Policy Optimization (MDPO);
- Different types of baselines:
  - Hindsight Credit Assignment;
  - Action-dependent baselines;
- IMPLEMENTATION MATTERS IN DEEP POLICY GRADIENTS: A CASE STUDY ON PPO AND TRPO

<https://arxiv.org/abs/2005.12729>





# Recap for RL

- What is RL?
- Markov Decision Process;
- Why don't we use it everywhere and where it is useful;
- V-function, Q-function;
- Value Iteration, Policy iteration;
- Least-Squared Value Iteration;
- Exploration-Exploitation trade-off;
- Experience Replay (Replay Buffer);
- Deep Q-Network (DQN);
- Policy Gradient Theorem, log-derivative trick;
- REINFORCE;
- Actor-Critic algorithm and A2C;